


File Management —:



Boot Block : Occupies beginning of filesystem, typically first sector. Contains bootstrap code read into memory to boot (init, pid 0).

Super Block : Information about filesystem, free blocks, pointers to inode list.

Inode List : List of inodes, where inodes provide file specific information.

Data Blocks : Data of files.

- x **Inode :** Contains necessary info for processes accessing the file such as file ownership, access rights, location of file data & file size in the filesystem.
 - Exist in static copy form on disk (Secondary), from where kernel reads them into an in-core form for manipulation.
 - Disk Inodes contain :
 - 1) File owner identifier :

d	rwx	rwx	rwx
file type	owner perms	group perms	other perms
d = dir - = file			

 - Ownership divided between individual owner & group (set of users given access)
 - SU has access to all files
 - 2) File Type : regular, directory, FIFO, socket, etc.

- 3) File Access Perms —
 - Owner, Group owner, Other users
 - Each class given access rights for read, write & execute
- 4) File Access Times —
 - Time of last modification (mtime), last access (atime) & last inode modification (ctime)
- 5) # Links to file —
 - # of symbolic (soft & hard) links to the file (aliases)
- 6) Table of Contents for disk addresses of file data
 - File stored in discontinuous disk blocks, accessed via kernel as logical byte stream. Blocks identified by inode data.
- 7) File Size —
 - File data addressable by # of bytes from beginning of file, starting from byte offset 0.
 - In-core Inodes contain, in addition to disk inodes file fields :
- 8) Status of inode —
 - Indication of properties :
 - x Inode is locked or not.
 - x A process is waiting for inode to be unlocked
 - x In-core rep. of inode is different as a result of change of inode data.
 - x In-core file rep. differs from disk as a result of change to file data.
 - x File is a mount point.
- 9) Logical device # of filesystem containing file
- 10) Inode # : Stored as linear array on disk, so kernel identifies disk inode by position.
- 11) Pointers to other incore inodes :
- 12) Reference Count : # instances of file active,
of processes accessing instance.

- x Kernel identifies inodes based on file system & inode # & allocates in-core inodes.
- x Maps device & inode # into a hash queue & searches queue.
- x On unsuccessful search, a new inode is allocated from a maintained free list.

$$\text{block \#} = \left\lfloor \frac{(\text{inode \#} - 1)}{\text{\# inodes per blk.}} \right\rfloor + \text{start } \text{block} \text{ address of inode list.}$$

(inode # = 1-based indexing)

$$\text{byte offset} = \left[(\text{inode \#} - 1) \% \text{\# inodes per blk.} \right] \times \text{Size of disk inode.}$$

(Indicates byte offset of inode in block)

- x In case of no inodes of free list, a process is terminated with an error instead of sleep to prevent starvation.

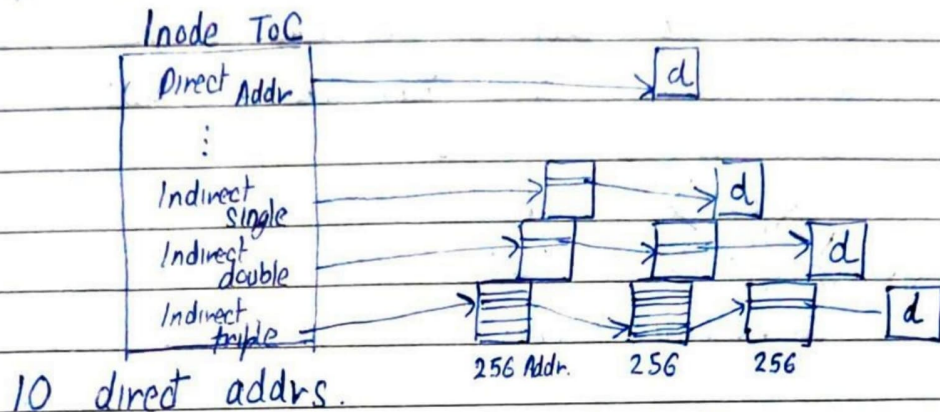
Releasing Inodes :

- On release, kernel decrements in-core reference count.
- When ref. count hits 0, if dirty write (in-core copy different from disk copy) then kernel initializes async write.
- Inode is then added to free list.
- If link count is 0, disk blocks for inode is freed, followed by freeing of inode.

Structure of Regular File - :

- x Inode's table of contents contain block numbers containing file data in discontinuous blocks.

- × Prefer discontinuous block-based allocation for efficiency & reduced fragmentation. One block at a time is allocated.
- × Table of Contents Structure :



For 1Ki blocks : (1Ki=1024)

10 direct blocks = 10KiB

1 single indirect = 256KiB

1 double indirect = 64 MiB

1 triple indirect = 16 GiB (16,777,216 B)

Addr = 0 : Indicates absence of data block.

Eg Let size of block = 1024 B

To access byte offset 9000, process needs to access

block address = $\lceil 9000 / 1024 \rceil \approx 9$ from ToC

leading to data block = ToC[9] = 367 (say)

To access byte offset 350000, process needs to access

block address = $\lceil 350000 / 1024 \rceil = 342$.

which lies in double indirect block 0 at address 76

so required data block = $ABlock_0[ToC[12]] = ABlock_0[76]$

⊗ Block address = 1 + Block offset in list.

Directories —

- x Provides hierarchical structure to filesystem.
- x Aids in conversion of file name to inode #.
- x Initial access : open, chdir, link syscalls.
- x Inode of '/' = 0.

/ <path> = absolute path

~~<./././~~ <path> = relative path

directory perms —

- r : Search dir.
- w : Write dirs & files.
- x : Open, access dir.

Eg: Inodes in '/home' :

Byte Offset	Inode #	File
0	12	etc.
16	16	..
32	162	etc

Eg: accessing /a/b/c

Fetch inode for /, given permission is available.
 Search for inode # of a, given it exists.
 Fetch inode for /a, given " " "
 Search for inode # of b, " " "
 Fetch inode of /a/b, " " "
 Search for inode of c, " " "
 Fetch inode of /a/b/c

- x Inode of CWD stored in U-Area of a process.

E.g. Fetch /etc/passwd

- x Fetch inode with # 0 (inode for /), (working dir = /)
- x Search inode # for 'etc' in inode of /, inode# = 83
past reading inode of / & its disk blocks.
- x 'etc' present, set working dir = /etc (inode# = 83)
- x Fetch inode with # 83 (inode for /etc) (working dir = /etc)
- x Search inode # for 'passwd' in inode of /etc, inode# = 2114
- x 'passwd' present, set working path = /etc/passwd (inode# = 2114)
- x No more path components, return inode# = 2114.

Super Block -

- x Holds information about :
 - x Size of filesystem
 - x Free Blocks (count)
 - x List of Free Blocks (available)
 - x Index of next free blocks list's free block
 - x Size of inode list
 - x # of free inodes
 - x List of usable free ~~blocks~~ inodes
 - x Index of next free inode
 - x Flag for information about modification
 - x Lock fields.
- x Fetching free inode for assignment to file (new file)
 - Inefficient way : Search ^{disk}inode list for free ^{disk}inode (type=0) & allocate (linear search, expensive)
 - Improvisation : Cache free inode #s in super block.
- x Remembered inodes :
 - When superblock is free (empty free inode list), kernel searches disk & places as many free inodes as possible on free inode list.


- When kernel needs more free inodes, search is resumed from last inode index that had been searched instead of the beginning. (remembered inode).
- Remembered Inode: Last inode from free inode list of super block. (Marked to trigger search).

Freeing Disk Inodes — (ifree)

- × On freeing of inode, add to superblock based on inode #: inode # lower than remembered inode implies update superblock free inode list & update remembered inode #. Inode # higher than remembered inode implies free inode not added to free inode list.
- × Free inode list is unsorted

Allocating Disk Blocks —

- × Free disk blocks cached in superblock for allocation to files by kernel for processes.
- × Data blocks stored in linked list, where each node contains array of free block #'s & one entry points to index of next disk block node on list.



Linked list structure shifted (elements copied to parent) when last block of superblock list is to be allocated.

- × Different structure for efficiency (discontiguous memory & large # of disk blocks & requirement of multiple blocks).