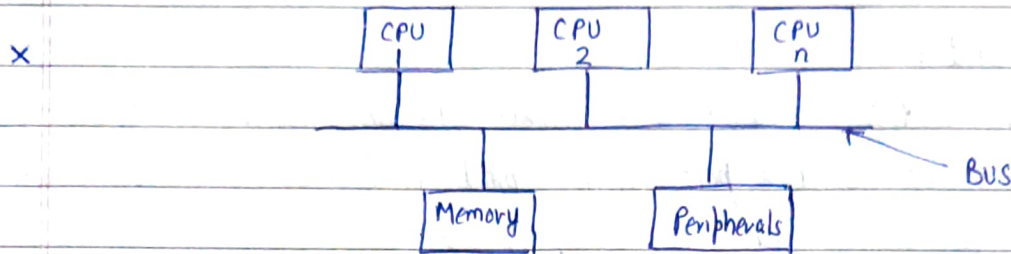


Multiprocessor Systems —



- Multiple CPUs within single PC, sharing System Bus, Memory & I/O devices
- CPUs execute independently, but exec. one copy of kernel.
- Processes behave as on uniprocessor system — semantics of syscalls are same, but processes allowed migration b/w processors.
- Allows processes to work in parallel, increasing throughput, reliability & cost-saving behavior (less memory cost).

x Problems :

- Inconsistency due to simultaneous execution in different processors.
- Solutions :
 - Execute critical actions on one processor (master & slave arch.)
 - Serialize access to critical regions w/ lock. primitives (semaphores)
 - Redesign algorithms to avoid critical sections.

x Master & Slave Processors :

- Master processor handles syscalls & interrupts.
- Slave processors execute in user mode & notifies master of syscalls.
- Process table now has field to indicate processor ID to exec. on (eg. master / slave).
- Scheduling algorithm now checks if process may be run on slave for slave processors & assigns highest priority process to each processor based on status (master or slave)
- Syscall invocation now marks processor ID to that of master if executed on slave & replace process in ready queue.

- Once process finishes syscall, reset processor ID to any (slave or master). Also, if master processor is required for another process, perform context switch to assign this process to a slave.
- Clock intr handler periodically reschedules processes to prevent monopolization.
- Challenges - :
 - Scheduler algorithm does not protect against same process being selected for execution on multiple processors.
 - Load Balancing : One processor may have high # of processes assigned, while others are idle.
- Solutions - :
 - Master processor may specify slave processor for process execution, thereby acting as scheduler.
 - Further, master permits more than one process to be assigned to a processor.
 - Master kernel distributes process load b/w processors.

x Semaphores w/ Multiprocessors :

- Partition kernel into critical regions such that atmost one processor executes in critical region at a time.
- Issues : Implementation of semaphores
Definition of critical regions.
- Sleep-Lock mechanism :

while lock set :

sleep (until lock is free);

set lock;

———— lock ————

free lock

wakeup processes sleeping on condition lock is set.

———— unlock ————

- Does not work on multiprocessor systems, as lock checked simultaneously by processors imply multiple processors lock a limited resource, leading to corruption due to incorrect lock state.

- Semaphore Init: true value

P: decrement (wait if value < 0)

V: increment (signal to waiting processes)

CP: Conditional P (decrement, returns true if value > 0 ,

If value ≤ 0 , unchange value, sleep until > 0)

- PPrim, VPrim: Processor locking primitives.

- Implementable using:

- Software instructions (array of semaphore values)

- HW instructions: read & clear (atomic)

read & clear: read value at location, clear value & set condition code as per original value.

- Lock: Control access to semaphore

- Value: Determine if process has access to critical region.

- Queue: Processes sleeping on semaphore.

- Semaphore now contains components: lock, value & queue.

- Problems: Infinite loops prevent preemption.

Solutions: Updated algorithms for P & V using PPrim & VPrim.

x P: lock semaphore lock using PPrim, next decrement semaphore value. If value is still ≥ 0 , release lock (VPrim) & return. Else, attempt sleep: check signals, if signal to interrupt sleep, increment semaphore value & if signal caught in kernel, release lock in VPrim & return -1 (abnormal wakeup in kernel mode). Else release lock & longjmp for context switch. If no signals, enqueue process in semaphore queue, release lock (VPrim) perform context switch & check signals before return.

x V: lock semaphore (PPrim), increment value. If value is still 0 or < 0 , dequeue from semaphore queue & make it ready to run (wakeup process). Lastly, release lock (VPrim).

while (value(semaphore) < 0) V(semaphore);

- Not equivalent to wakeup, as in wakeup, signals are sent to indicate to processes that resource is available, but V for multiprocessor directly assigns resource, leading to inconsistency. (non atomic)
- Conditional P: provide solution to possible deadlock situations.
If CP is false, release held resources to prevent deadlock.

x Types of Multiprocessing —:

- Asymmetric (Master & Slave)
- Symmetric (Peer-Peer Relation)

- ASMP Challenges :

- Performance Degradation (Master becomes Bottleneck)

- SMP :

- Every processor has shared copy of OS.
- On syscall, CPU receiving syscall traps to kernel to execute.
- May pick same process for execution, preventable using mutexes.
- Allow modularity in OS. (each processor runs non interacting part of OS).
- Challenges :
 - Complex OS
 - Large main memory requirement.